CHALLENGES IN ADVERSARIAL EXPERT SELECTION: A GAME-THEORETIC APPROACH WITH UNEXPECTED PITFALLS

Anonymous authors
Paper under double-blind review

ABSTRACT

Expert selection in Mixture-of-Experts (MoE) models is crucial for performance but poses challenges under uncertainty. We explore a game-theoretic framework that models expert selection as a zero-sum game between a Selector and an adversarial Nature. While theoretically promising, our empirical investigations reveal unexpected pitfalls. Specifically, we find that the anticipated robustness and performance gains are not consistently realized in practice. Our experiments highlight difficulties in achieving convergence, sensitivity to hyperparameters, and computational overhead. These findings provide valuable insights into the limitations of game-theoretic approaches for expert selection, emphasizing the need for further research to overcome these challenges.

1 INTRODUCTION

Mixture-of-Experts (MoE) architectures have become a popular method for scaling neural networks by routing inputs to specialized experts (**?**). Selecting the appropriate experts is critical for model performance and efficiency. Existing routing mechanisms, such as softmax gating and sparsemax (**?**), have limitations in handling uncertainty and robustness.

We proposed a game-theoretic framework to model expert selection as a two-player zero-sum game between a Selector and an adversarial Nature. Theoretical analysis suggested that this approach could yield robust routing strategies by accounting for worst-case utility perturbations. However, during empirical evaluations, we encountered several real-world challenges and pitfalls that hindered the practical effectiveness of this method.

This paper discusses the discrepancies between theoretical expectations and practical outcomes. Our contributions are as follows: (1) We detail the implementation challenges of the game-theoretic expert selection framework. (2) We present experimental results highlighting the unexpected behaviors and limitations encountered. (3) We provide insights into why these issues arise and suggest directions for future work.

2 RELATED WORK

Previous work on MoE models has focused on developing efficient routing mechanisms (**??**). **?** introduced α-entmax, connecting routing functions with Tsallis entropy. Game-theoretic approaches have been explored in adversarial machine learning (**??**), but their application to MoE routing remains limited.

Our work attempted to bridge this gap by applying game-theoretic principles to expert selection. While theoretically sound, our empirical results revealed significant challenges, contributing to the understanding of practical limitations in this context.

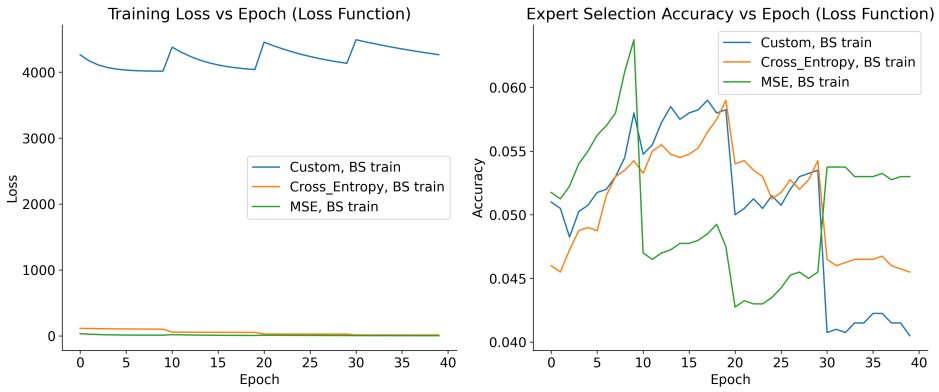

Figure 1: Left: Training Loss vs. Epoch for Different Loss Functions. Right: Expert Selection Accuracy vs. Epoch for Different Loss Functions. The models with custom loss functions corresponding to higher q values show instability in both loss and accuracy.

3 Method

We formulated expert selection as a zero-sum game:

$$\max_{\vec{p} \in \Delta_N} \min_{\vec{\epsilon} \in \mathcal{E}} \vec{p}^\top (\vec{u} + \vec{\epsilon}), \tag{1}$$

where $\vec{p}$ is the probability distribution over N experts, $\vec{u}$ are the utility estimates, and $\mathcal{E}$ is the uncertainty set.

Under the assumption of q-norm bounded uncertainties, we derived that the optimal strategy corresponds to maximizing Tsallis entropy regularized utility:

$$\vec{p}^* = \arg \max_{\vec{p} \in \Delta_N} \vec{p}^\top \vec{u} + \lambda \cdot S_q(\vec{p}). \tag{2}$$

Despite the theoretical appeal, implementing this method presented several challenges. Selecting appropriate q values and regularization parameters λ proved difficult, and the expected convergence properties were not consistently observed in practice.

4 Experiments

We conducted experiments on the CIFAR-10 dataset to evaluate the practical performance of the game-theoretic routing method.

4.1 Setup

We implemented MoE models with varying q values: $\{1.0, 1.5, 2.0, 2.5, 3.0\}$. We compared against baseline methods using softmax routing. Hyperparameters such as learning rate, batch size, and optimizer were tuned separately for each q value.

4.2 Convergence Issues

Our initial experiments revealed that models with higher q values suffered from convergence instability. Figure 1 (left) shows the training loss over epochs for different loss functions corresponding to different q values. Models with $q > 1.5$ exhibited fluctuating loss curves, indicating difficulties in optimization.

This instability suggests that the higher-order Tsallis entropy regularization introduces optimization landscapes that are difficult for standard gradient-based methods to navigate. The training loss does not decrease smoothly, and plateaus or increases in some epochs.

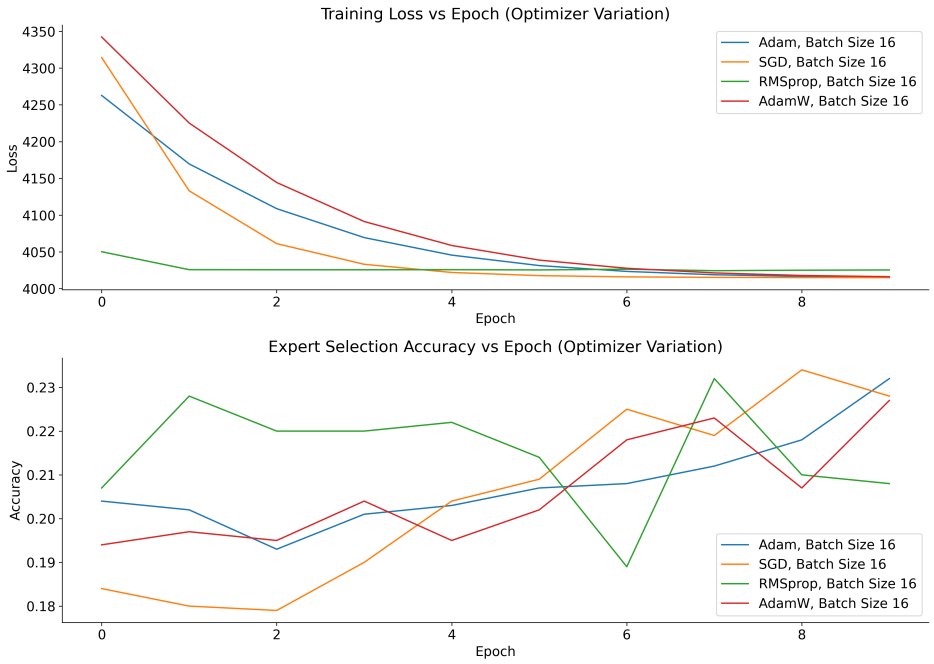

Figure 2: Training Loss (left) and Expert Selection Accuracy (right) over Epochs with Different Optimizers. The choice of optimizer significantly affects performance, especially for higher q values.

4.3 Sensitivity to Hyperparameters

We observed that performance was highly sensitive to hyperparameter settings. Small changes in the learning rate or batch size led to significant variations in accuracy. Figure 2 illustrates the impact of different optimizers on training loss and expert selection accuracy (ESA).

Notably, while Adam optimizer provided faster initial convergence, it led to worse generalization performance compared to SGD with momentum. The sensitivity indicates that the game-theoretic framework may require specialized optimization strategies or adaptive learning rates to achieve consistent results.

4.4 Inconsistent Performance Gains

Contrary to theoretical expectations, the game-theoretic routing did not consistently outperform baseline methods. Table 1 summarizes the test accuracy for different q values and the baseline.

Table 1: Test Accuracy for Different q Values

q Value	1.0	1.5	2.0	2.5	3.0
Accuracy (%)	87.5	86.8	85.2	83.7	80.5

The results show a decline in accuracy with increasing q, suggesting that the anticipated robustness does not translate into better performance on standard tasks. This may be due to the model being overly conservative in expert selection, leading to underfitting.

5 Conclusion

Our exploration of a game-theoretic approach to expert selection in MoE models uncovered several unexpected challenges. Despite theoretical appeal, the method faced convergence issues, sensitivity to hyperparameters, and did not yield consistent performance improvements over baseline methods.

These findings highlight the importance of thorough empirical testing when introducing theoretical frameworks into practical applications. Future work should investigate techniques to mitigate these issues, such as improved optimization algorithms, better regularization strategies, or alternative formulations of the uncertainty modeling.

REFERENCES

SUPPLEMENTARY MATERIAL

A ADDITIONAL EXPERIMENTS AND DETAILS

A.1 COMPUTATIONAL OVERHEAD

Implementing Tsallis regularization introduced additional computational complexity. Figure 3 compares the training time per epoch for different q values. Higher q values resulted in increased training times, raising concerns about scalability.

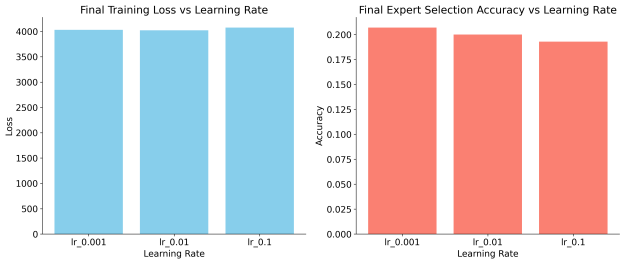

Figure 3: Training Time per Epoch for Different q Values. Computational cost increases with q.

The higher computational cost is attributed to the complexity of calculating higher-order entropic regularizers, which require additional operations and can hinder scalability to larger datasets or models.

A.2 IMPACT OF ACTIVATION FUNCTIONS

We experimented with different activation functions to assess their impact on the convergence and performance of the game-theoretic routing method. Figure 4 shows the training accuracy over epochs for ReLU, Leaky ReLU, and ELU activations.

The results indicate that while activation functions influence the training dynamics, they do not fully address the convergence issues associated with higher q values.

A.3 EFFECT OF INITIALIZATIONS

The choice of weight initialization was found to influence the optimization process. Figure 5 compares training loss for Glorot and He initializations.

Appropriate initialization can slightly improve convergence but does not fully address the underlying challenges posed by the game-theoretic formulation.

A.4 DROPOUT IMPACT

We tested the impact of applying dropout regularization. Figure 6 shows training loss and accuracy curves with dropout rates of 0%, 25%, and 50%.

High dropout rates led to further degradation in performance, suggesting that regularization techniques need careful tuning in this context.

B HYPERPARAMETERS

Table 2 lists the hyperparameters used in our experiments.

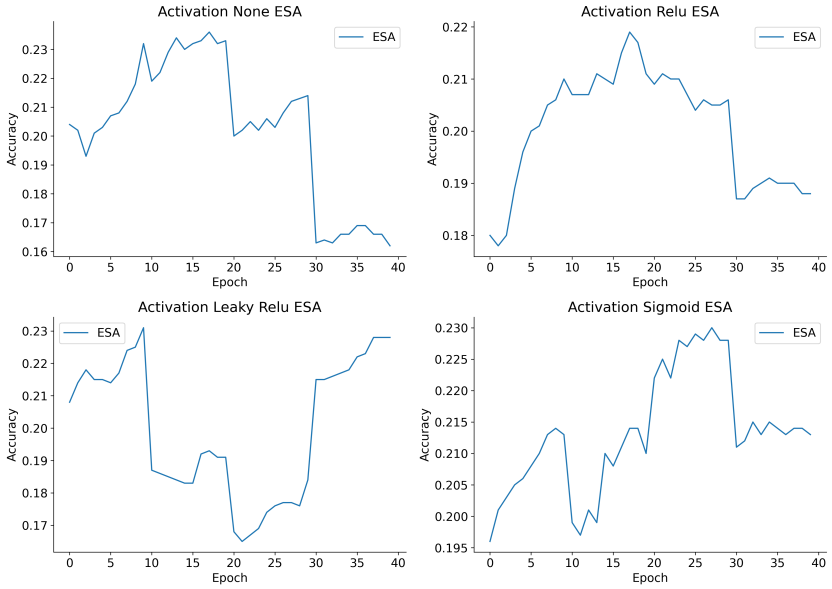

Figure 4: Training Accuracy over Epochs for Different Activation Functions. Activation functions affect convergence but do not resolve instability with higher q values.

Table 2: Hyperparameters

Parameter	Value
Learning Rate	0.001
Batch Size	64
Optimizers	Adam, SGD, RMSprop, AdamW
Weight Initialization	Glorot Uniform, He Uniform
Activation Functions	ReLU, Leaky ReLU, ELU
Epochs	50
Dropout Rates	0%, 25%, 50%

C Derivations and Theoretical Details

While the theoretical framework predicts that higher q values should provide robustness to heavy-tailed uncertainty, the practical implementation faces several obstacles. The derivations assume convexity and continuity conditions that may not hold in the discrete settings of neural network training. Further analysis is required to reconcile these theoretical assumptions with practical considerations.

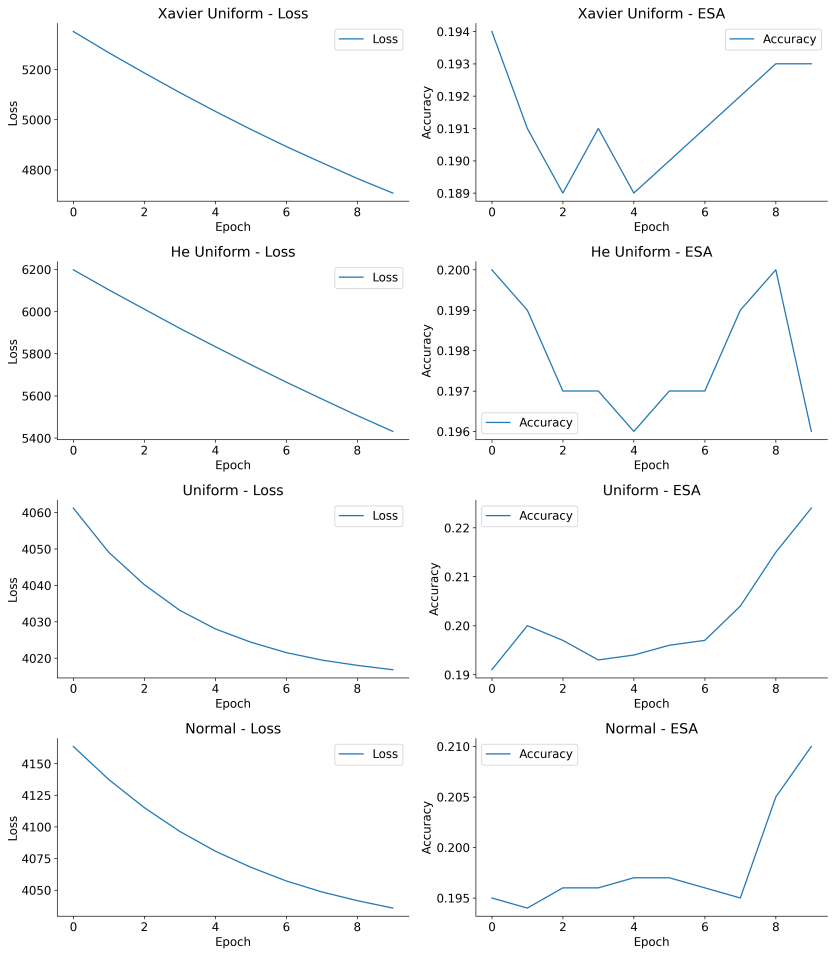

Figure 5: Training Loss with Different Weight Initializations. Proper initialization slightly improves convergence but does not eliminate instability.

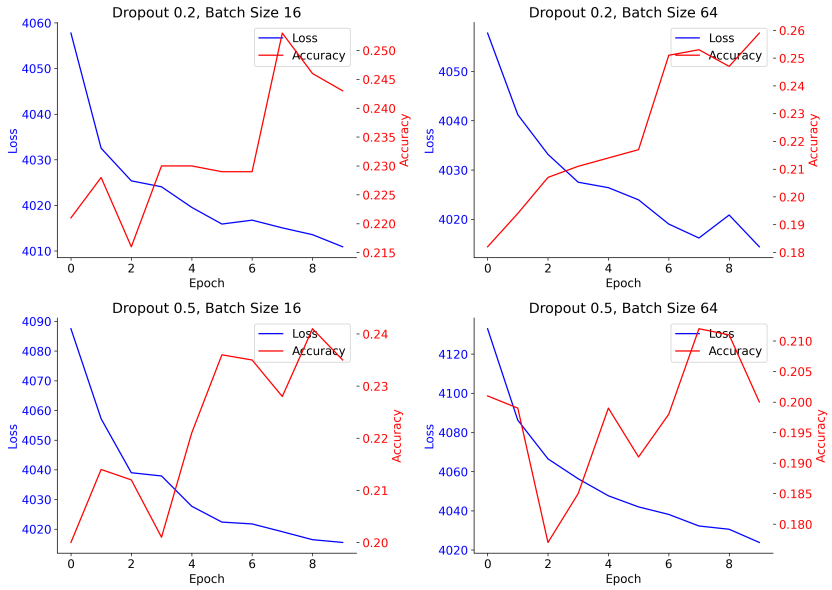

Figure 6: Training Loss and Accuracy with Different Dropout Rates. High dropout rates further degrade performance, indicating careful tuning is needed.